

Data structures can typically be stored either in one solid block (contiguous) or in many small blocks, each with a memory address for the next block (linked). What are some pros and cons of each approach? Contiguous pros: good physical locality of memory means faster in practice. No wasted space for memory addresses. Linked pros: infinite adding so long as there are open spaces big enough for element + address.

Your application



Operating System



Memory

1	2	3	4
5	6	7	8
9	10	11	12

Your application

Operating System

Memory



1	2	3	4
5	6	7	8
9	10	11	12

I just need one
block of memory for
now...doesn't matter
where!

Your application




Okay, I'm adding a new item. I'll use the previous cell you gave me to point to the next

Operating System



Memory



LINKED LISTS



The Linked List DS

- Data structure used for *list*, *stack*, *queue*, *deque* ADTs etc.
- Uses *nodes* which encapsulate a *value* and pointer(s)
- Main entity holds reference(s) to just a head and/or tail node
 - the "*handle(s)*"
- Each node then *points* to the *next* and/or *previous* node
 - "*singly-linked*" (unidirectional) vs. "*doubly-linked*" (bidirectional)





Linked List

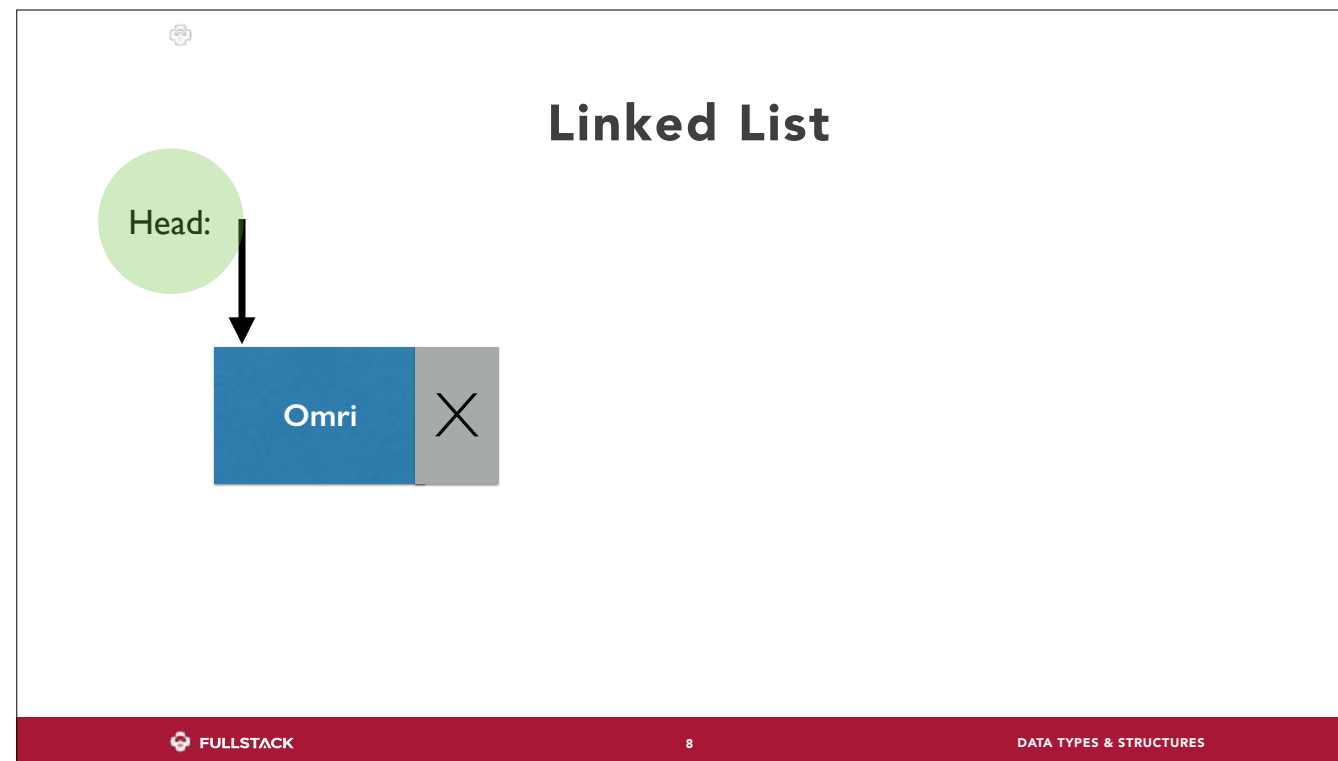
Head:

 FULLSTACK

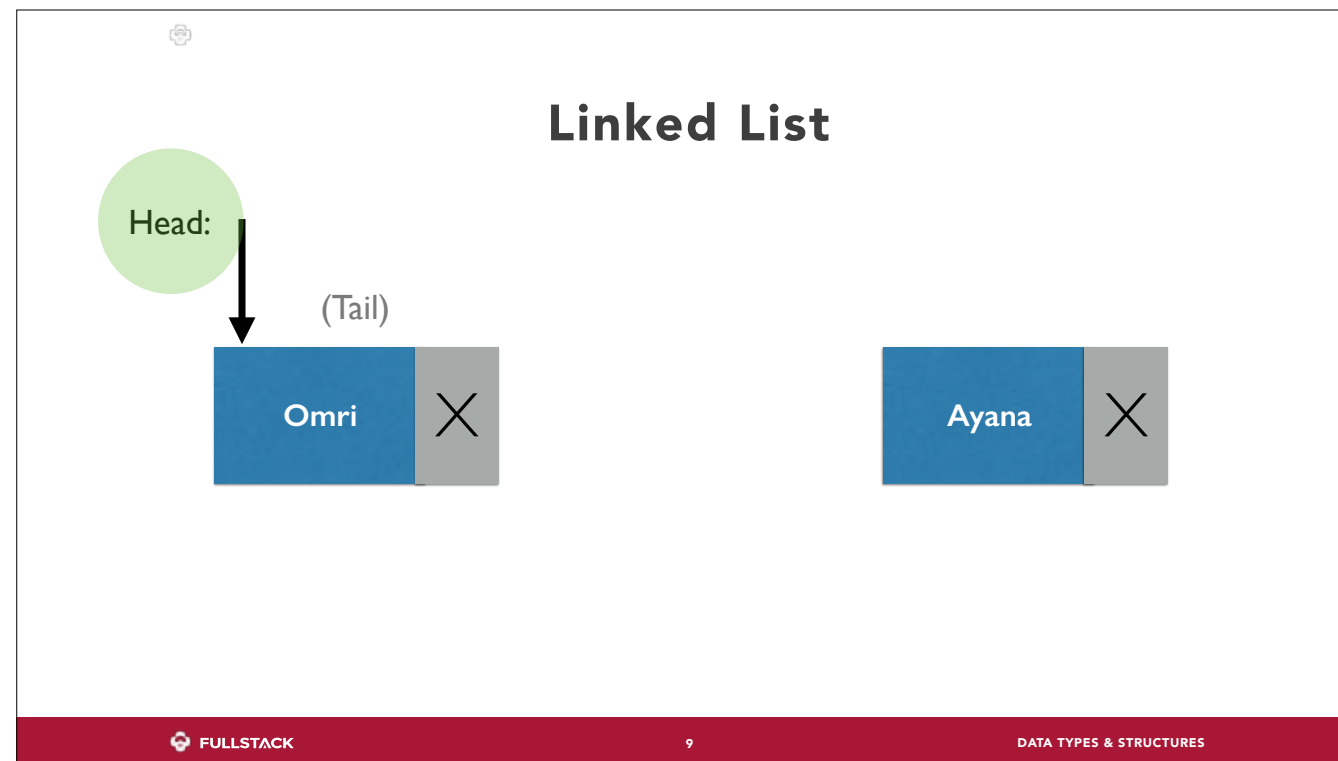
7

DATA TYPES & STRUCTURES

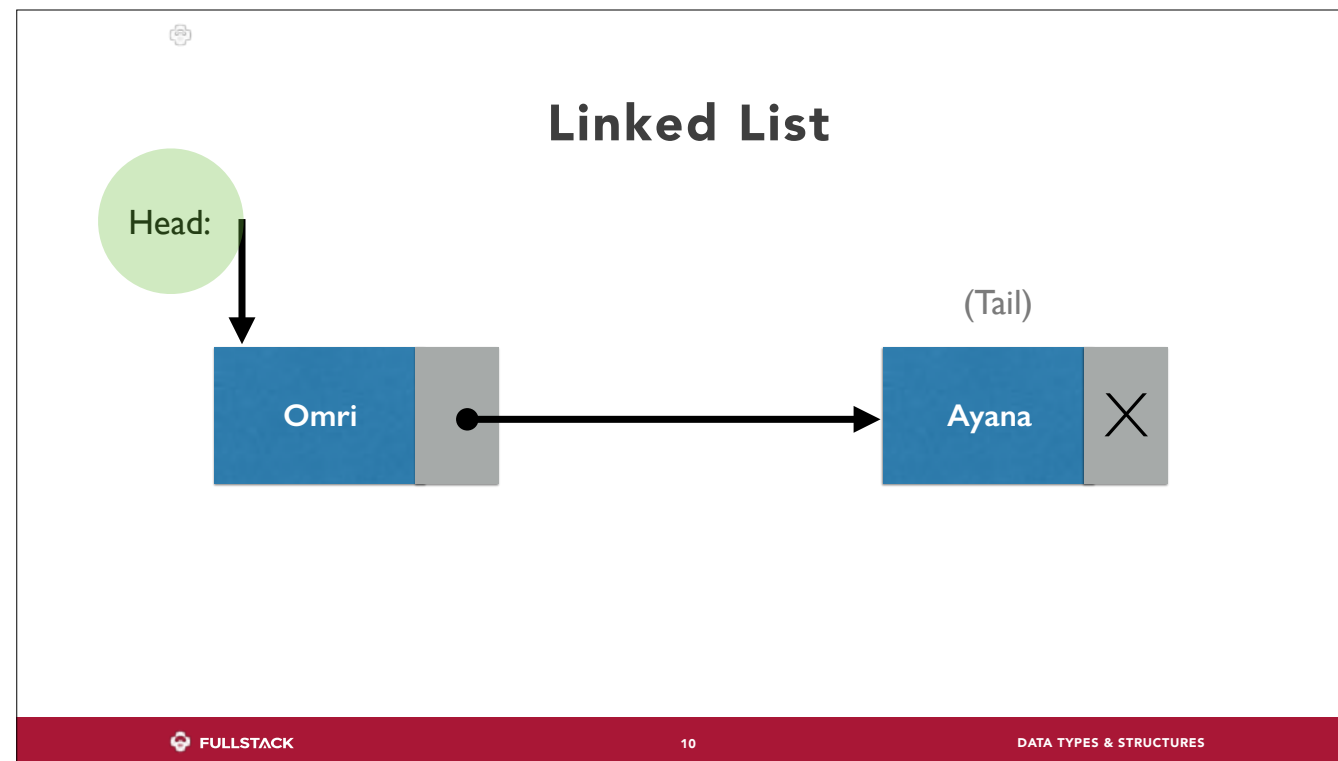
When it comes to implementations, the LinkedList is often considered synonymous with its handle or handles (in this case just an empty "head" reference).



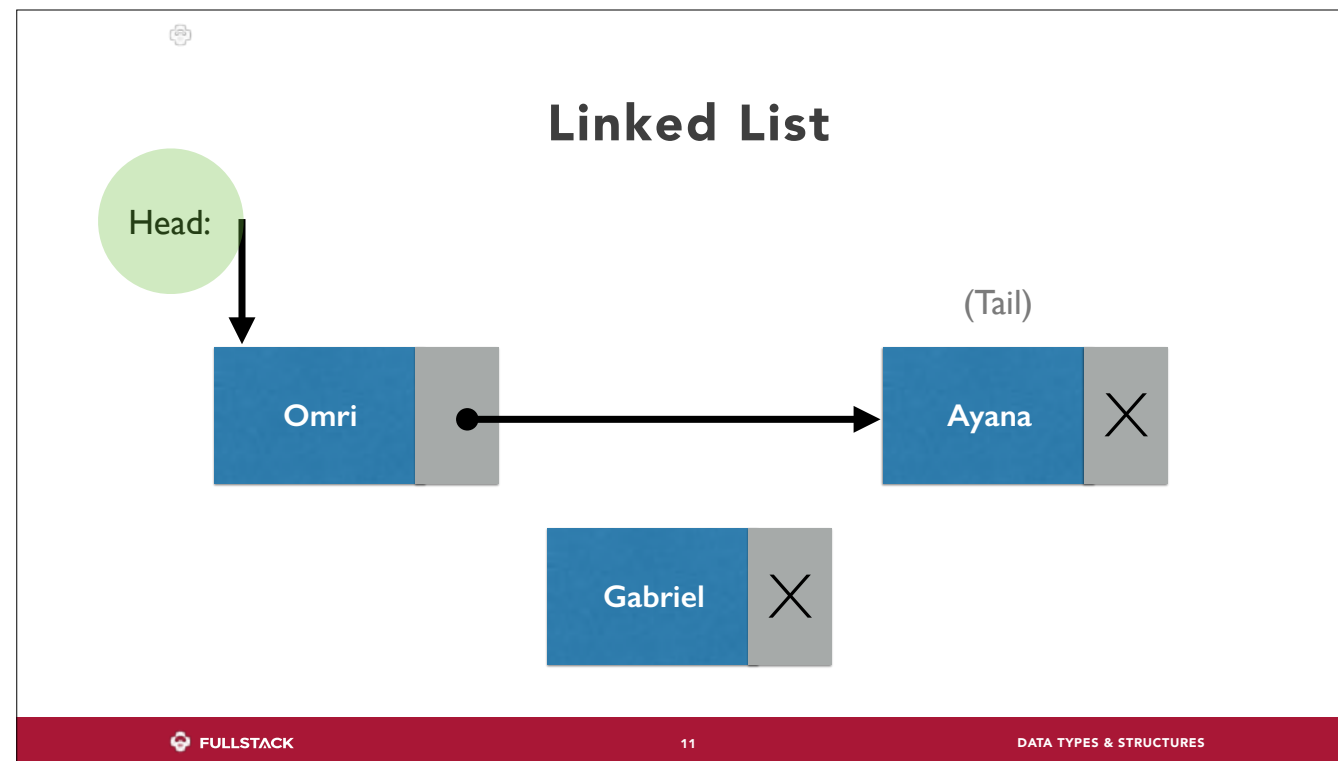
This is a Linked List containing one node. The node has a value (Omri) and an empty `next` pointer. In this case it only has one pointer (no `previous`) so this is going to be a singly-linked list.



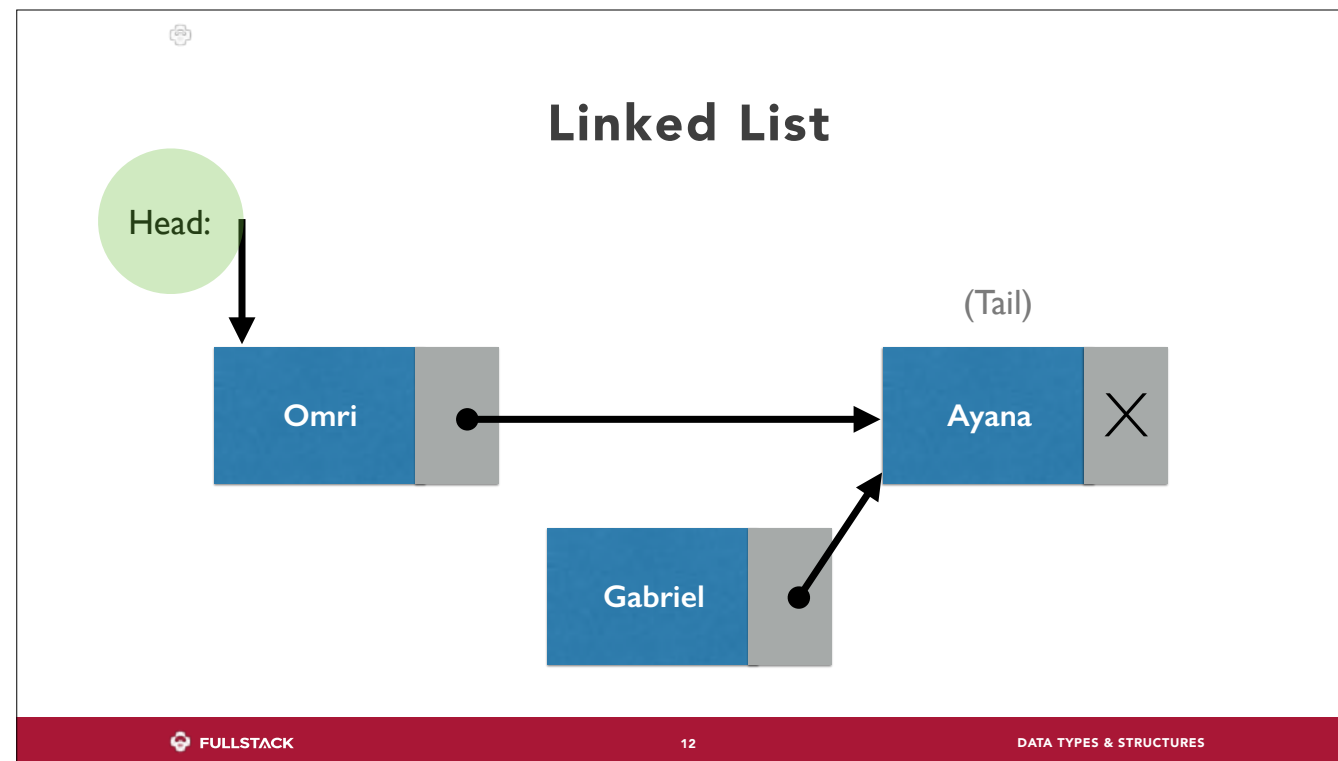
Partway through adding to the tail — we have created a new Ayana node, but it isn't connected to our list yet.



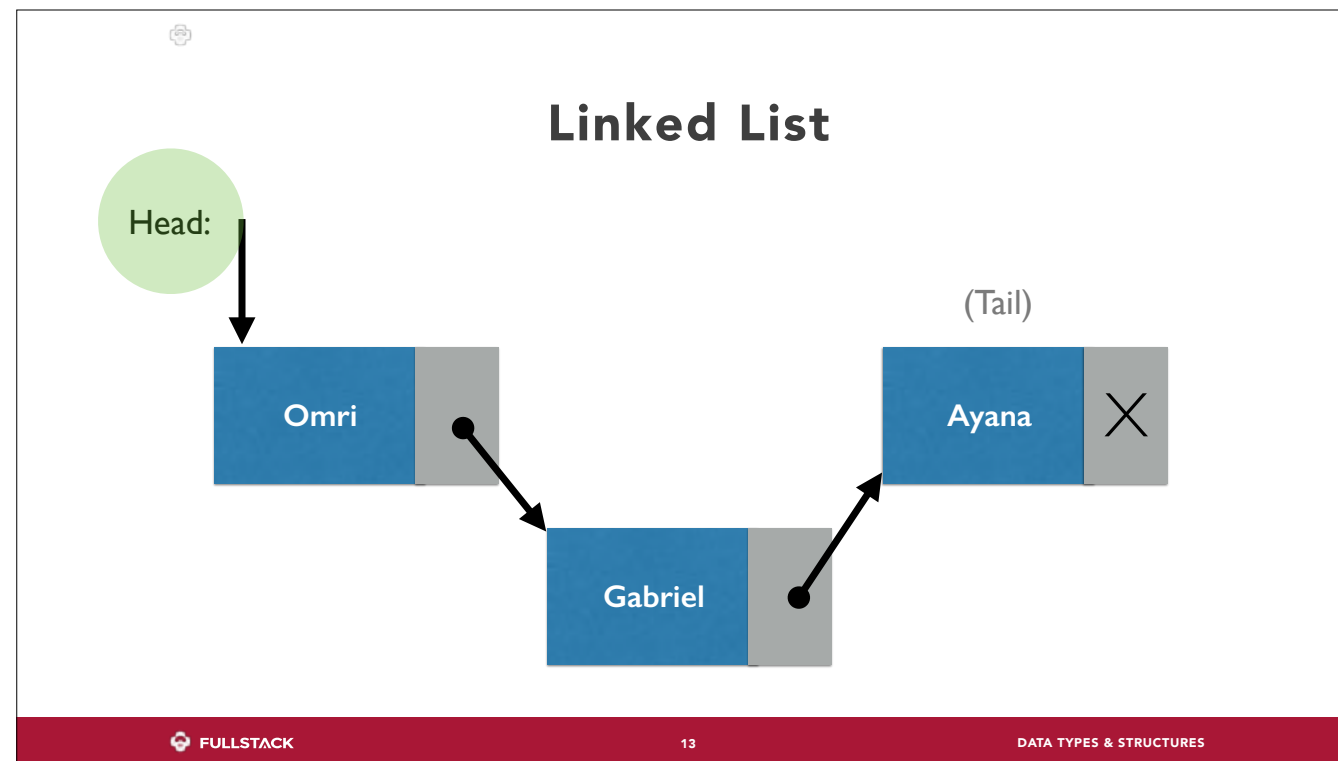
Set the Omri node's `next` pointer to the Ayana node, and we would say that Ayana is the new tail node. Can I get to Omri from Ayana? If I want to reach Ayana, I have to start at `head` and then follow `next`.



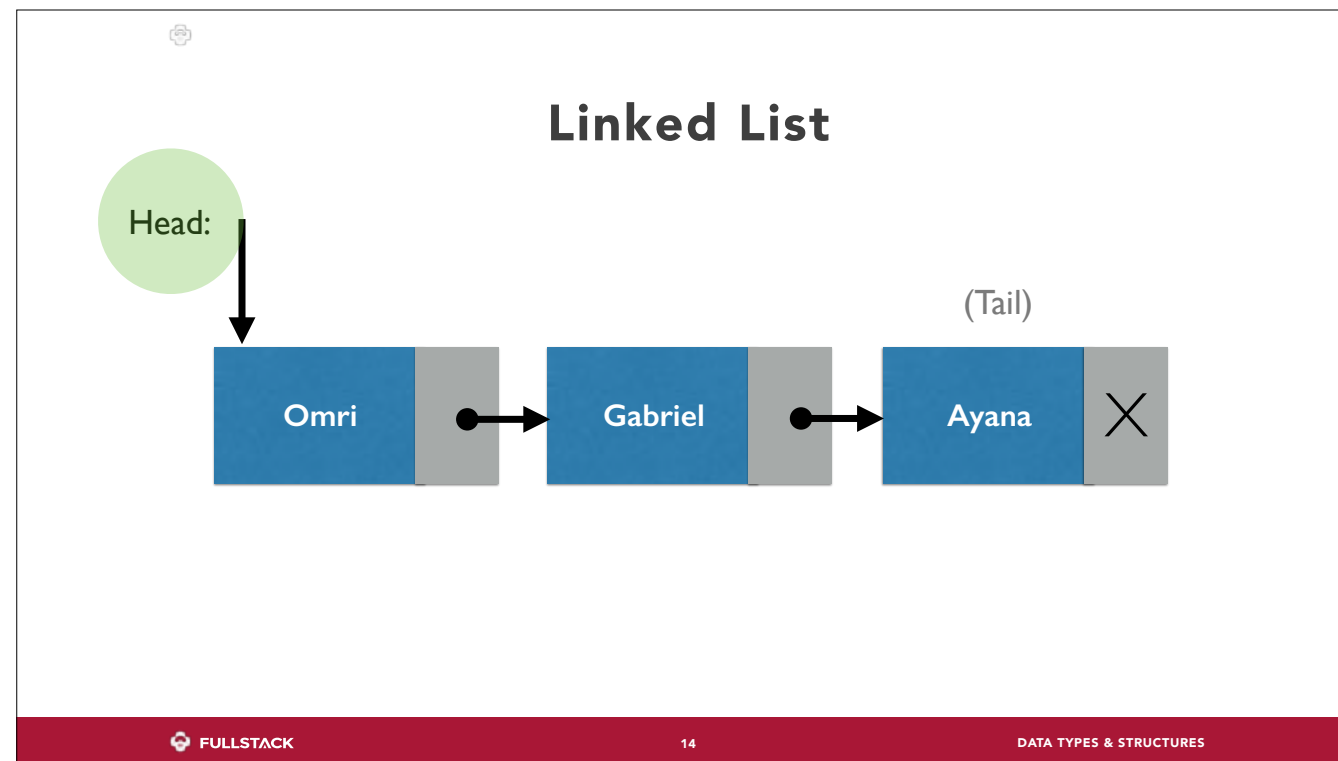
What if I want to add something in the middle of the list?



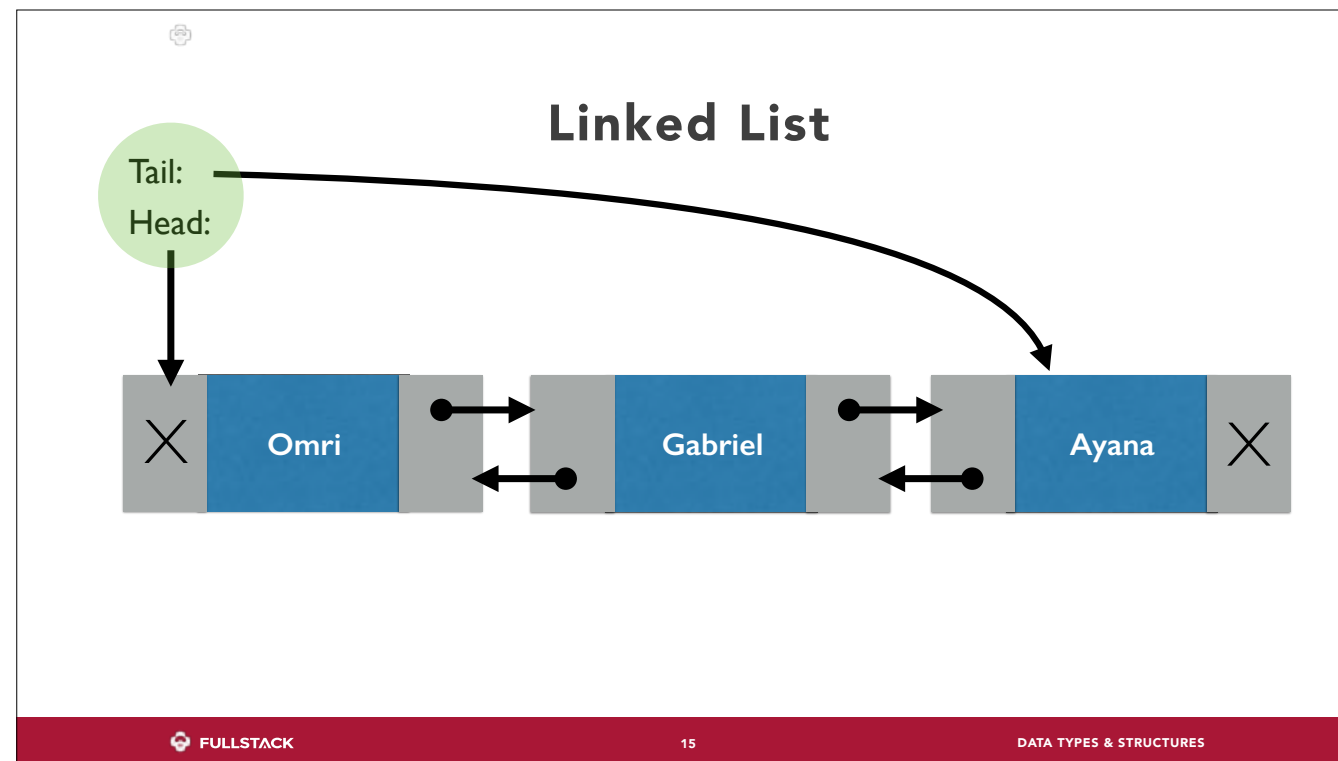
The best thing about Linked Lists is that *insertion* (adding in the middle) always takes the same number of steps, regardless of how big the list is. Again, we call that "constant time." First, point the new node's `next` to the old later node.



Then, change the previous node's `next` to point to the new node.



That's it!



Here is an example of a doubly-linked list with head and tail pointers. Think about the pros and cons of this approach.



(some) pros/cons: Linked Lists vs Arrays

Operation	Linked List	Typed Array
Reach element in middle	Must crawl through nodes	Constant time
Insert in middle or start	Constant time (if we have ref).	Must move all following elements
Add element to end	With handle, constant time	Constant time
Space per element	Container + element + pointer(s)	Just element!
Total space	Grows as needed	Pre-reserved & limited*
Physical locality	Not likely	Best possible

**Dynamic* arrays keep track of their limit, and if it is exceeded, the array is copied into a new bigger array (which is expensive but hopefully happens infrequently).



Part 1 of the workshop!