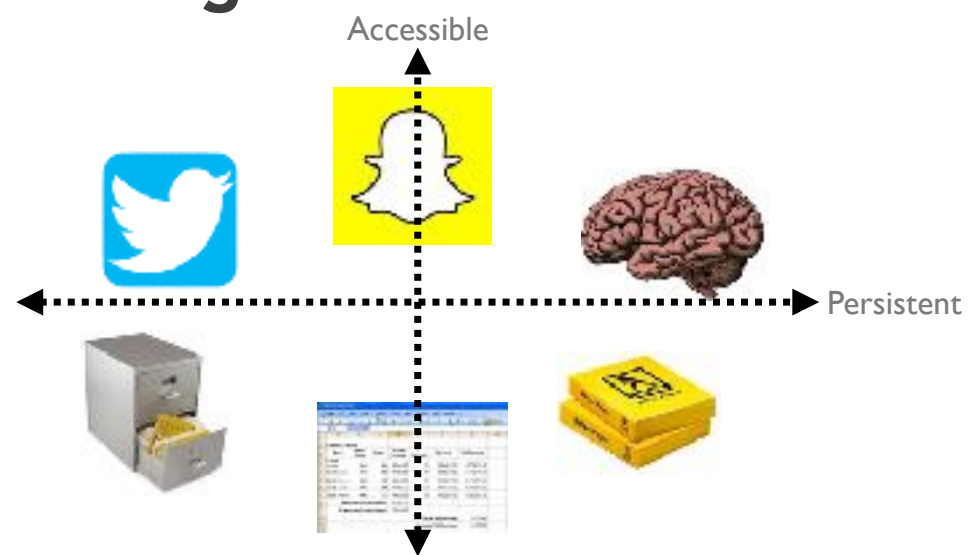


Intro to Databases

SQL

What is a database?

Things that hold info



A database **persists** information
and is **accessible** via code

```
graph TD; A[accessible] --- B[organized]; A --- C[queryable]; A --- D[manageable]
```

organized **queryable** **manageable**

Organized: Standard Storage Formatting

- DBs are a collection of **Tables** (or *relations*)
- Tables have **Columns** (*attributes* / *fields*) that describe **Rows** (*instances* / *tuples*)
- Duplicate rows are not allowed
- Rows often have a **primary key** (unique identifier)

Imagine an excel spreadsheet.
natural way to express information.
Objects often nicely map to tables

Table / Relation

Column / Attribute / Field

Column / Attribute / Field

Column / Attribute / Field

	ID	Name	Type
Row / Tuple / Instance	1	Pikachu	lightning
Row / Tuple / Instance	2	Squirtle	water
Row / Tuple / Instance	3	Charmander	fire
Row / Tuple / Instance	4	Bulbasaur	grass

Queryable: via a Standard Language

- A simple, structured query language: SQL
- Declarative (vs. imperative)
- No more hand-rolled algorithms / data structures
- DBMS picks an efficient execution strategy based on indexes, data, workload etc.



SQL

```
-- Pikachu, I choose you!  
SELECT id, name  
  FROM pokemon  
 WHERE type = 'lightning'  
LIMIT 1
```


Manageable: Easy, Safe, Performant

- Offloads work and requisite understanding of programming
- Knowledge is portable
- Abstraction
- Transfer data between systems
- DBMS can make certain guarantees
 - prevent unsafe operations
 - built-in redundancies
 - handle multiple users, threads

ACID Guarantees

- **Atomicity**
- **Consistency**
- **Isolation**
- **Durability**

Atomicity - A set of database operations that must occur together. All or nothing. A transaction cannot partially finish, it must either fail or complete.

Consistency - If a process has a writer, no other process can read from it, and no other process can write to it, which gives us consistent information at all times.

Isolation - Multiple clients can make queries to read and update without the risk of deadlock or starvation.

Durability - Store information without power (flash drive / hard drive, for example doesn't need power and still stores data)

Atomic Transactions

- **atomic transaction: A set of database operations that must occur together**
 - i.e. A debit to one bank account, and a credit to another
- **A transaction must either succeed or fail; it cannot partially complete.**
- **Every database query is represented by a transaction**

We might also say that a database is atomic or a database has the property of atomicity.

We can define where transactions start and end. You may need to as a programmer

Consistency

- **Specify rules that columns need to follow**
 - Gender column can only contain M, F, or U.
 - Savings account must start with S or checking with C
 - Column cannot be null
- **Protect the database from inconsistencies and simplify software logic**
 - Allows software to make assumptions about underlying data

Resource Management

- Processes can be readers and writers
- Files can have many readers
- If a process has a writer, no other process can read from it, and no other process can write to it

Proposed File Scheme

- Suppose that we have decided not to use a database and instead store our data in a series of files.
- How might our setup fail to serve queries from multiple users?

Deadlock



- Three files: A, B, and C
- Process 1 needs files A then B for writing
- Process 2 needs files B then C for writing
- Process 3 needs files C then A for writing
- What happens if all three processes start a database request at once?

Databases give us concurrency (Isolation)

- Multiple clients can make queries to read and update without the risk of deadlock or starvation.

Persistence/Durability

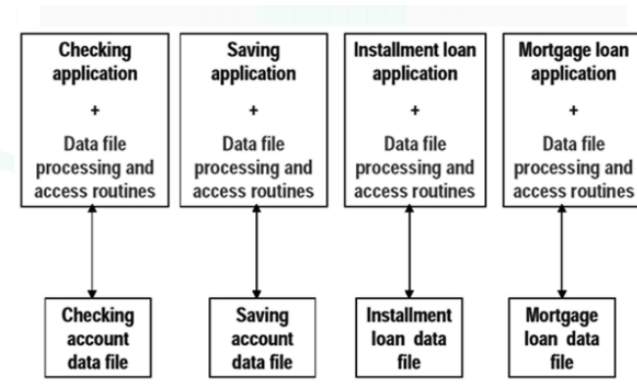
- Files are also persistence (store information without power)

- What happens when two people try to access the same file at the same time?
- the second process waits and retries until the first process finishes?

How Did We End Up Here?

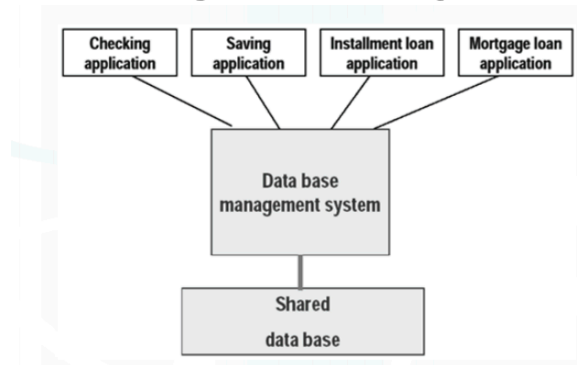
Before Relational DBs (ca. < 1970s)

- Data stored in custom “data files”
- Queried via application-specific code
- Advantages
 - Middle layer not needed
 - Solutions customized for each application
- Disadvantages
 - Hard to change the system
 - Knowledge not compounding
 - Data-transfer is difficult



- Separate file for every topic. Routines and functions for dealing with each particular file.
- But every company formatted their files differently based on developer whims and needed their own accessor functions.
- Making changes was really hard seeing as formats were inflexible and etched in stone.

Database Management Systems (DBMS)



- One layer and language to store and access data
- Sold as a way for “non-technical people” to manage data

Ubiquitous
Standardized
Don't need to be a programmer
common way of thinking about data. And a common language for querying it.

“Future users of large data banks must be protected from having to know how the data is organized in the machine (the internal representation).”

– E. F. CODD,
A RELATIONAL MODEL OF DATA FOR
LARGE SHARED DATA BANKS

Relational Databases & Logic



- 1969: Edgar Frank "Ted" Codd outlines *relational model* of data
- Wrote Alpha (never implemented) as a *query language*
- IBM slow to adopt his ideas
 - Competitors started to do so
 - IBM team formed without Codd, created **Structured English Query Lang**
- **SEQUEL** way better than what came before
 - 1979: copied by Larry Ellison (from pre-launch papers / talks!) as "SQL"
- **SQL became the standard (ANSI 1986, ISO 1987)**
 - Codd continued to fault SQL compared to his theoretical model
 - The Third Manifesto: solve the *object-relational impedance mismatch*

Appreciating Databases

- Ubiquitous
- Standardized
- Complex / deep
- Powerful: database admins are
 - Feared by developers
 - ...but also taken for granted until things break
 - Befriended by business people
 - Contacted by the government for secret data (e.g. NSA)

Progression of Databases

- **Navigational (< 1970s)**
 - More common during tape era; entries had references to next entries.
- **Relational (> 1970s)**
 - Based on relational (table-based) logic, see E.F. Codd.
- **NoSQL (> 2000s)**
 - "Not only SQL" — document storage, for example.

Navigational: common during tape era, entries had references to the next entries.

Relational: based on relational (table-based) logic, EF Codd.

NoSQL ("not only SQL"): document storage.

RDBMS vs NoSQL

- **A DBMS doesn't *have* to be relational**
 - Remember, DBMS is just an application that intelligently stores data and can answer requests to manage that data
- **Lately, many "NoSQL" or non-relational DBMSs have been gaining popularity**
 - Graph databases (e.g. Neo4j)
 - Document databases (e.g. MongoDB)
 - Hybrids (e.g. PostgreSQL)
- **RDBMSs still remain the #1 DB option for now**



Some well-known rDBMSs





Why PostgreSQL?

- Advanced, powerful, and popular
- Rapid open source development
- Highly extensible (stored procedures)
- Deep SQL standards compliance
- NoSQL ("Not Only SQL"), objective support
- Excellent transactions / ACID reliability; focus on integrity
- Multi-user management / administration

History of PostgreSQL

- 1970s at UC Berkeley:
Interactive **G**raphics **R**etrieval **S**ystem (INGRES)
- 1980s: POSTGRES ("Post-Ingres")
- 1995: POSTQUEL and Postgres95.
 - monitor -> psql
- 1996: Adopted by the open source community
 - Ongoing: stability, testing, documentation, new features
 - PostgreSQL

psql



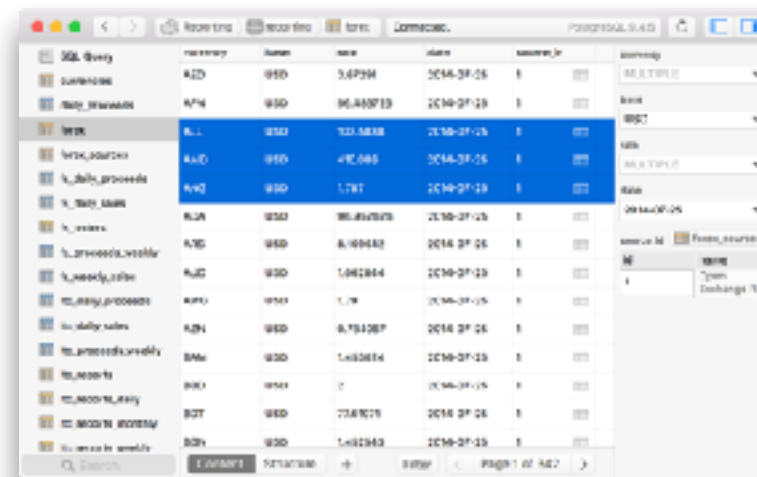
The screenshot shows a psql terminal window with the following content:

```
psql (13.1)
Type in help to get help.
13.1=# \l
          List of databases
  Name                | Owner  | Encoding | Collate | Ctype  | Access privileges
-----+-----+-----+-----+-----+-----
template0             | postgres | UTF8     | en_US.UTF-8 | en_US.UTF-8 | 
template1             | postgres | UTF8     | en_US.UTF-8 | en_US.UTF-8 | 
postgres              | postgres | UTF8     | en_US.UTF-8 | en_US.UTF-8 | 
postgres1             | postgres | UTF8     | en_US.UTF-8 | en_US.UTF-8 | 
postgres2             | postgres | UTF8     | en_US.UTF-8 | en_US.UTF-8 | 
postgres3             | postgres | UTF8     | en_US.UTF-8 | en_US.UTF-8 | 
postgres4             | postgres | UTF8     | en_US.UTF-8 | en_US.UTF-8 | 
postgres5             | postgres | UTF8     | en_US.UTF-8 | en_US.UTF-8 | 
postgres6             | postgres | UTF8     | en_US.UTF-8 | en_US.UTF-8 | 
postgres7             | postgres | UTF8     | en_US.UTF-8 | en_US.UTF-8 | 
postgres8             | postgres | UTF8     | en_US.UTF-8 | en_US.UTF-8 | 
postgres9             | postgres | UTF8     | en_US.UTF-8 | en_US.UTF-8 | 
postgres10            | postgres | UTF8     | en_US.UTF-8 | en_US.UTF-8 | 
postgres11            | postgres | UTF8     | en_US.UTF-8 | en_US.UTF-8 | 
postgres12            | postgres | UTF8     | en_US.UTF-8 | en_US.UTF-8 | 
postgres13            | postgres | UTF8     | en_US.UTF-8 | en_US.UTF-8 | 
postgres14            | postgres | UTF8     | en_US.UTF-8 | en_US.UTF-8 | 
postgres15            | postgres | UTF8     | en_US.UTF-8 | en_US.UTF-8 | 
postgres16            | postgres | UTF8     | en_US.UTF-8 | en_US.UTF-8 | 
postgres17            | postgres | UTF8     | en_US.UTF-8 | en_US.UTF-8 | 
postgres18            | postgres | UTF8     | en_US.UTF-8 | en_US.UTF-8 | 
postgres19            | postgres | UTF8     | en_US.UTF-8 | en_US.UTF-8 | 
postgres20            | postgres | UTF8     | en_US.UTF-8 | en_US.UTF-8 | 
(11 rows)

13.1=# \d
          List of relations
 Schema | Name          | Type  | Temp |
-----+-----+-----+-----+
 public | pg_class      | class |      |
 public | pg_indexes    | index |      |
(2 rows)
```



Postico



The screenshot shows the Postico application interface. On the left, there's a sidebar with a tree view of the database structure. The main area displays a table with the following columns: name, id, amount, date, and status. The table contains 15 rows of data. The first row is highlighted in blue. On the right side, there's a panel with various settings and filters, including a dropdown menu for 'MULTIPLE' and a search bar.

name	id	amount	date	status
daily_sales	1	100.00	2016-07-25	1
daily_sales	2	200.00	2016-07-25	1
daily_sales	3	300.00	2016-07-25	1
daily_sales	4	400.00	2016-07-25	1
daily_sales	5	500.00	2016-07-25	1
daily_sales	6	600.00	2016-07-25	1
daily_sales	7	700.00	2016-07-25	1
daily_sales	8	800.00	2016-07-25	1
daily_sales	9	900.00	2016-07-25	1
daily_sales	10	1000.00	2016-07-25	1
daily_sales	11	1100.00	2016-07-25	1
daily_sales	12	1200.00	2016-07-25	1
daily_sales	13	1300.00	2016-07-25	1
daily_sales	14	1400.00	2016-07-25	1
daily_sales	15	1500.00	2016-07-25	1

LAB